

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [2]: from sklearn.linear_model import LogisticRegression
```

```
In [3]: df=pd.read_csv("C9_Data.csv")
df
```

```
Out[3]:
```

	row_id	user_id	timestamp	gate_id
0	0	18	2022-07-29 09:08:54	7
1	1	18	2022-07-29 09:09:54	9
2	2	18	2022-07-29 09:09:54	9
3	3	18	2022-07-29 09:10:06	5
4	4	18	2022-07-29 09:10:08	5
...
37513	37513	6	2022-12-31 20:38:56	11
37514	37514	6	2022-12-31 20:39:22	6
37515	37515	6	2022-12-31 20:39:23	6
37516	37516	6	2022-12-31 20:39:31	9
37517	37517	6	2022-12-31 20:39:31	9

37518 rows × 4 columns

```
In [4]: df=df.dropna()
df
```

```
Out[4]:
```

	row_id	user_id	timestamp	gate_id
0	0	18	2022-07-29 09:08:54	7
1	1	18	2022-07-29 09:09:54	9
2	2	18	2022-07-29 09:09:54	9
3	3	18	2022-07-29 09:10:06	5
4	4	18	2022-07-29 09:10:08	5
...
37513	37513	6	2022-12-31 20:38:56	11
37514	37514	6	2022-12-31 20:39:22	6
37515	37515	6	2022-12-31 20:39:23	6

	row_id	user_id	timestamp	gate_id
37516	37516	6	2022-12-31 20:39:31	9
37517	37517	6	2022-12-31 20:39:31	9

37518 rows × 4 columns

In [5]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 37518 entries, 0 to 37517
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   row_id      37518 non-null  int64
1   user_id     37518 non-null  int64
2   timestamp   37518 non-null  object
3   gate_id     37518 non-null  int64
dtypes: int64(3), object(1)
memory usage: 1.4+ MB
```

In [6]: `df.columns`

Out[6]: Index(['row_id', 'user_id', 'timestamp', 'gate_id'], dtype='object')

In [7]: `feature_matrix=df[['row_id', 'user_id']]`
`target_vector=df['gate_id']`

In [8]: `feature_matrix.shape`

Out[8]: (37518, 2)

In [9]: `target_vector.shape`

Out[9]: (37518,)

In [10]: `from sklearn.preprocessing import StandardScaler`

In [11]: `fs=StandardScaler().fit_transform(feature_matrix)`

In [12]: `logr=LogisticRegression()`
`logr.fit(fs,target_vector)`

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:763: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(
Out[12]: LogisticRegression()

In [13]: observation=[[1,2]]

In [14]: prediction=logr.predict(observation)
print(prediction)

[3]

In [15]: logr.classes_

Out[15]: array([-1, 0, 1, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16],
dtype=int64)

In [16]: logr.predict_proba(observation)[0][0]

Out[16]: 0.005365176788164149

In [17]: logr.predict_proba(observation)

Out[17]: array([[5.36517679e-03, 2.43221075e-05, 9.36568351e-05, 2.22025633e-01,
2.19695882e-01, 7.52352405e-02, 5.84513730e-02, 7.17956781e-02,
2.68284044e-03, 7.98655513e-02, 1.24425419e-01, 1.07054385e-01,
2.51118120e-03, 7.57336969e-03, 2.68214159e-05, 2.29125763e-02,
2.60893089e-04]])

In [18]: x=df[['row_id', 'user_id']]
y=df['gate_id']

In [19]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)

In [20]: from sklearn.linear_model import LinearRegression
lr=LinearRegression()
lr.fit(x_train,y_train)

Out[20]: LinearRegression()

In [21]: lr.intercept_

Out[21]: 7.278299996623737


In [22]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
coeff

Out[22]:

	Co-efficient
row_id	-0.000005
user_id	-0.012829

```
In [23]: prediction =lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[23]: <matplotlib.collections.PathCollection at 0x2099ed6c4c0>



```
In [24]: lr.score(x_test,y_test)
```

Out[24]: 0.005925010775539419

```
In [25]: lr.score(x_train,y_train)
```

Out[25]: 0.005346973822092593